


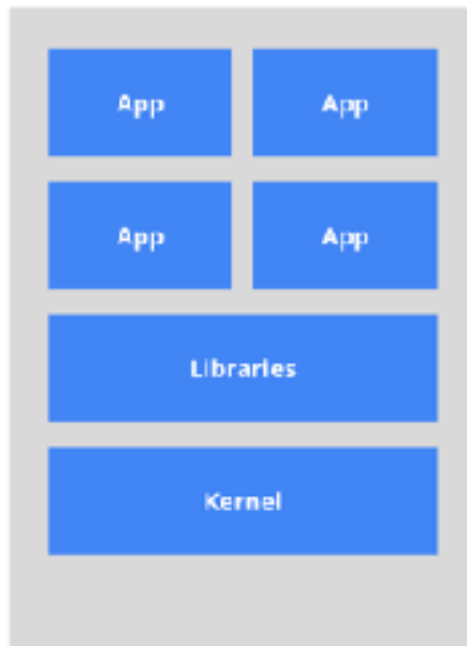
Docker & K8S

Thibault SAUSSAC




What's containers ?

The old way: Applications on host



*Heavyweight, non-portable
Relies on OS package manager*

The new way: Deploy containers



*Small and fast, portable
Uses OS-level virtualization*




What's Docker ? (vs virtual machines)

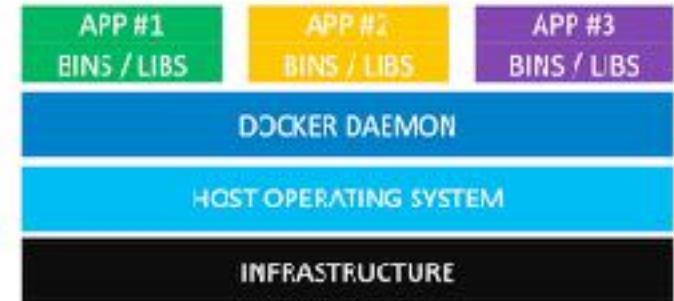
Les conteneurs Docker peuvent être considérés comme des VM sans leur propre système d'exploitation

Ils partagent tous le système d'exploitation hôte et sont séparés à l'aide d'outils du kernel tels que cgroups, namespaces, ..

Du point de vue de conteneur, chaque conteneur a son propre système de fichiers.




Virtual Machines




Docker Containers

What's Docker ? (vs virtual machines)



Virtual Machines



Docker Containers



Feature	VM	Container
Virtualization Boundary	Lowest	Highest
Resource abstraction	Hypervisor Interface	System Interface
Boot time	Fast, seconds	Very fast, milliseconds
Performance	Some hypervisor overhead	Close to that of bare-metal server
Density	High, dozens of VMs per host	Much higher, six to eight times as many containers or VMs on the same hardware system
Efficiency, resource utilization	Efficient	Highly efficient
Isolation	Full isolation with guaranteed resources is possible	Only process-level isolation
Resource management	Resources are managed by the host OS and guest OS.	Resources are managed only by the host OS.
Storage	Persistent storage is supported. Virtual disks are presented as the OS drives.	Singlebox containers do not support persistent storage; see Table 2 and Section 5 for details.
Security	VMs offer the security of a dedicated operating system.	The attack surface is larger as all the containers on a host run on a single shared Linux kernel (see Section 7).
Flexibility	Good	Separating and abstracting the interfaces between the host system and container makes the coding much easier.
Portability	VMs are portable between systems running the same hypervisor.	A container can be moved across any Linux server that supports the container runtime environment (see Section 8.2).
Maturity	Management products for virtualized environments have evolved over the years to ensure enterprise-level security, reliability, scalability, and availability.	It is relatively new but its ecosystem is growing rapidly.

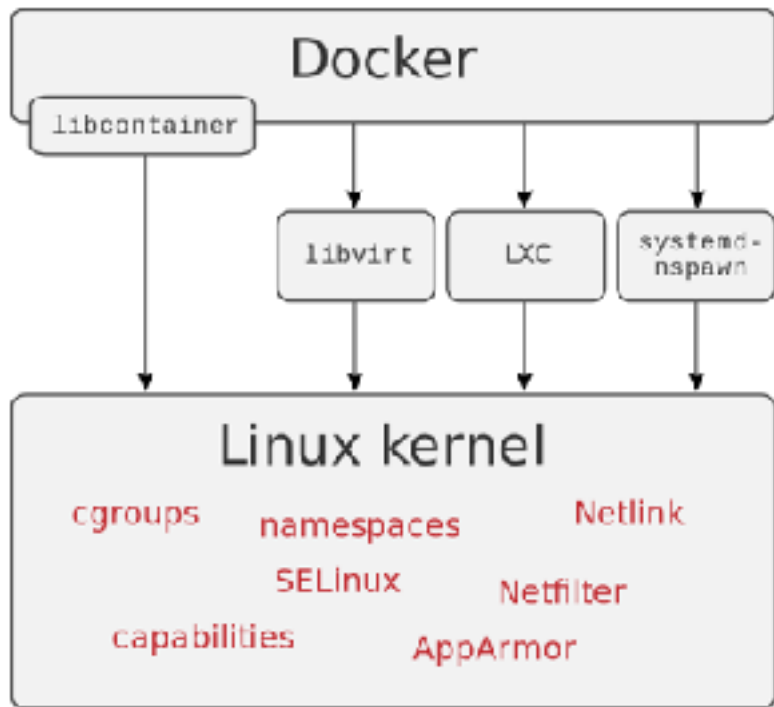
Table 1: Comparison of VMs and Container Features

Docker vs. Virtual machines

What's Docker ?



Les interfaces docker avec le kernel Linux



How it works ?



Dockerfile : La recette de cuisine

Un fichier qui décrit comment on build le conteneur

```
1  # base container
2  FROM python:3
3  # add file(s)
4  ADD my_script.py /
5  # command to run while building the container
6  RUN pip install pystrich
7  # command to run when container starts
8  CMD ["python", "/myscript.py"]
```

Et la commande, « `docker build` » pour build l'image

How it works ?



Dockerfile : La recette de cuisine

1 `# base container`

2 `FROM python:3`

3 `# add file(s)`

4 `ADD my_script.py /`

5 `# command to run while building the container`

6 `RUN pip install pystrich`

7 `# command to run when container starts`

8 `CMD ["python", "/myscript.py"]`

Start from an existing image

Ajouter un fichier dans à la racine

RUN une commande
(Installation des dep python)

La commande qu'on lance lorsque le
Conteneur démarre


How it works ?




Images vs Containers :

Image: fichier immuable contenant le code source, les bibliothèques, les dépendances, les outils et autres fichiers nécessaires à l'exécution d'une application.

Conteneur: environnement d'exécution virtualisé (en cours d'exécution)



How it works ?



How it works ?



```
$ docker
Commands:
  attach      Attach to a running container
  build       Build an image from a Dockerfile
  commit      Create a new image from a container's changes
  cp          Copy files/folders from a container's filesystem to the host path
  create      Create a new container
  diff        Inspect changes on a container's filesystem
  events      Get real time events from the server
  exec        Run a command in an existing container
  export      Stream the contents of a container as a tar archive
  history     Show the history of an image
  images      List images
  import      Create a new filesystem image from the contents of a tarball
  info        Display system-wide information
  inspect     Return low-level information on a container
  kill        Kill a running container
  load        Load an image from a tar archive
  login       Register or log in to a Docker registry server
```

```
logout      Log out from a Docker registry server
logs        Fetch the logs of a container
port        Lookup the public-facing port that is NAT-ed to PRIVATE_PORT
pause       Pause all processes within a container
ps          List containers
pull        Pull an image or a repository from a Docker registry server
push        Push an image or a repository to a Docker registry server
restart     Restart a running container
rm          Remove one or more containers
rmi         Remove one or more images
run         Run a command in a new container
save        Save an image to a tar archive
search      Search for an image on the Docker Hub
start       Start a stopped container
stop        Stop a running container
tag         Tag an image into a repository
top         Lookup the running processes of a container
unpause     Unpause a paused container
version     Show the Docker version information
wait        Block until a container stops, then print its exit code
```

How it works ?



Nb:

Le conteneur s'arrête si le processus s'arrête.

Si tu veux lancer une image Debian et rien d'autre il est necessaries d'explicité run bash :

```
docker run -it debian /bin/bash
```

(-i « run interactively » , -t « allocate a pseudo-tty »)

Une fois que le conteneur est arrêté, vous perdez les données à l'interieur.

En cas de besoin de persistance, 2 options : mount a host' folder dans le conteneur, ou ajouter un « Docker volume ».

The registry

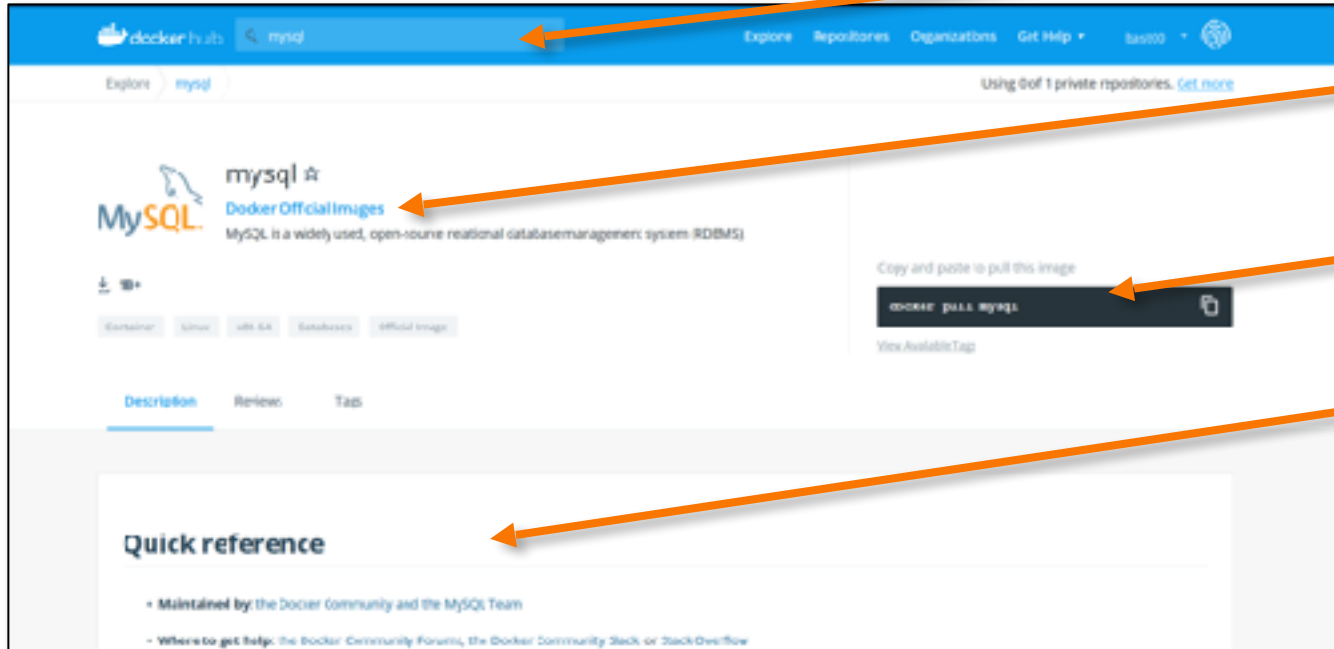
Retrieve images



`docker pull [imagename]`



The registry *hub.docker.com*



Search

Official image

Download the image

Documentation



Quick setup

Parametre parameters :

- name: nom du conteneur
- d: en background (detach)
- e: variable d'env

How to use this image

Start a mysql server instance

Starting a MySQL instance is simple:

```
$ docker run --name some-mysql --env MYSQL_ROOT_PASSWORD=secret-pw -d mysql:tag
```

... where `some-mysql` is the name you want to assign to your container, `secret-pw` is the password to be set for the MySQL root user and `tag` is the tag specifying the MySQL version you want. See the list above for relevant tags.

Connect to MySQL from the MySQL command line client

The following command starts another `mysql` container, attaches to it, and runs the `mysql` command line client against your `some-mysql` container. Of course you'll need to

The registry *hub.docker.com*



```
1 # Base container
2 FROM python:3
3 # Add file(s) source(host) dest(filesystem of the container)
4 ADD my_script.py /
5 # command executed while building the container
6 RUN pip install pystrich
7 # command executed when the container start
8 CMD [ "python", "./my_script.py" ]
```

Précédemment nous utilisions une image de base python:3.

Cela signifie qu'on utilise l'image docker avec le tag 3 ce qui correspond à « 3.8.5-buster »

L'image correspondante est automatiquement téléchargé depuis docker hub quand tu build ton image.

https://hub.docker.com/_/python

Shared Tags

- 3.9.0rc1, 3.9-rc, rc:
 - 3.9.0rc1-buster
 - 3.9.0rc1-windowsservercore-ltsc2016
 - 3.9.0rc1-windowsservercore-1809
- 3.9.0rc1-windowsservercore, 3.9-rc-windowsservercore, rc-windowsservercore:
 - 3.9.0rc1-windowsservercore-ltsc2016
 - 3.9.0rc1-windowsservercore-1809
- 3.8.5, 3.8, 3, latest:
 - 3.8.5-buster
 - 3.8.5-windowsservercore-ltsc2016
 - 3.8.5-windowsservercore-1809

The registry *hub.docker.com*

Push your own images :

FREE:

Unlimited public repos

One private repo

PRO (individuals): 5\$/months

Unlimited private repo


Pro tools for developers

TEAM: 7\$/user/month

Collaboration & Mgt tools

`docker pull [imagename]`

`docker push`



Prod : how to setup the prod application :



```
version: '3.0'

services:
  webserver:
    image: wordpress
    container_name: wp_web
    ports:
      - 8080:80
    links:
      - dbserver:mysql
    environment:
      WORDPRESS_DB_PASSWORD: 6zcznAEjLWp79P

  dbserver:
    image: mysql:latest
    container_name: wp_db
    environment:
      MYSQL_ROOT_PASSWORD: 6zcznAEjLWp79P
```

Le service s'appel « webserver »

L'image s'appel « wordpress »

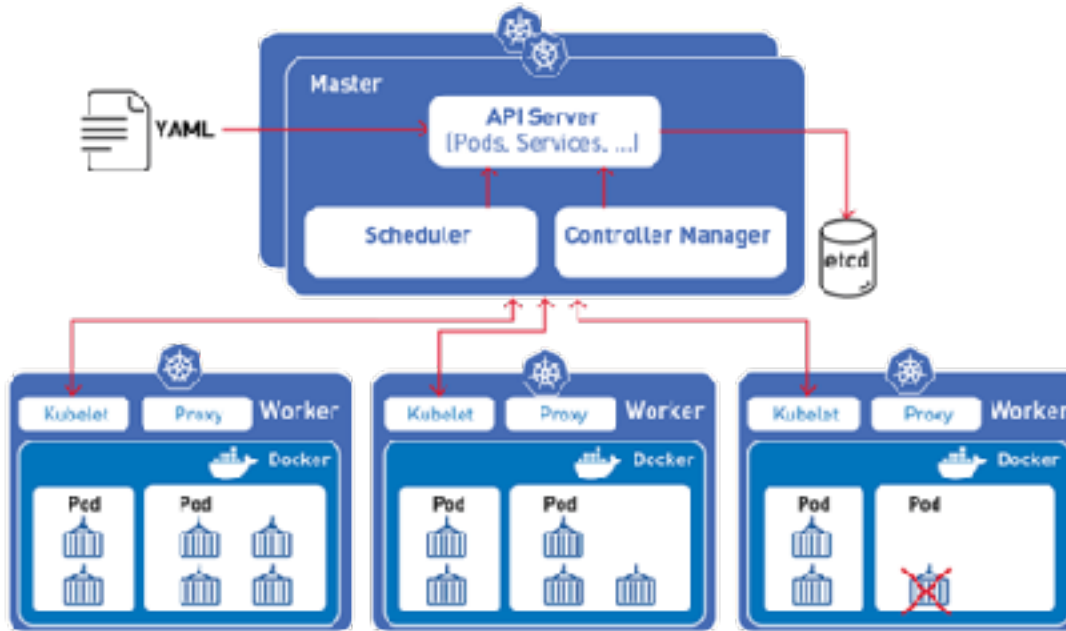
Map le port de l'host 8080 vers le port du conteneur 80

service dbserver ira chercher mysql

Environment variable

Pour utiliser ce fichier:
« docker-compose up »

Il ira automatiquement télécharger WP et MySQL



- Les hosts sont des « *nodes* »
- Groupe de conteneurs dans les « *pods* »
- Plusieurs pods par nodes
- Containers in pods :
 - Shared storage
 - Shared unique cluster IP address

Kubernetes




Master node, 3 processes:

- kube-apiserver
- kube-controller-manager
- kube-scheduler

Entrypoint / API

Assign pods to nodes

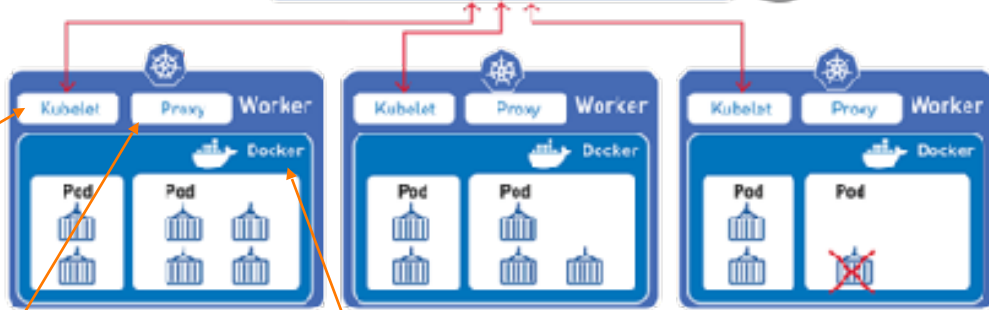
Control the system



Connect to master

Network services

Container runtime (e.g. Docker)



Nodes, 2 processes:

- kubelet
- kube-proxy & container runtime

Kubernetes : How to use



CLI : kubectl

Dashboard

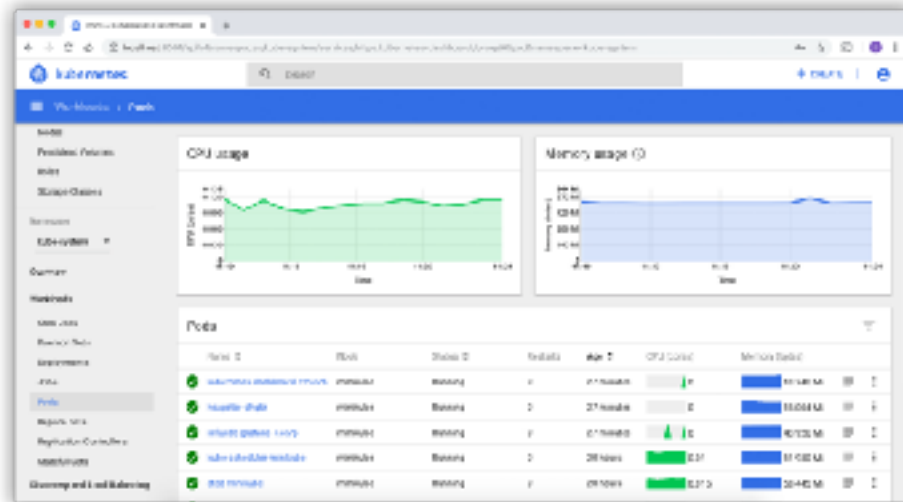

```
kubectl [action] [resource]
```

e.g.

```
kubectl get pods – list the pods
```

```
kubectl get nodes – list the nodes
```

```
kubectl create deployment mydeploy \  
--image=rabbitmq:latest
```



Créera un « *deployment* » qui appellera un conteneur basé sur RabbitMQ dispo sur Docker Hub.

Par défaut il crée 1 container dans 1 pod.

Kubernetes : How to use - Example



```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-765bf4c7b4-v6w0x	1/1	Running	0	10m

1 – list of pods on the cluster

Info : This container is the result of a deployment (here the deployment is « kubernetes-bootcamp »).

Kubernetes : How to use - Example



```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-765bf4c7b4-vvdx	1/1	Running	0	10m

1 – list of pods on the cluster

```
$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
kubernetes-bootcamp-765bf4c7b4	1	1	1	10m

2 – list of « *replicaset* »

The **replicaset** is a constraint applied to pods. If the number of pods don't match the desired number, it'll create or delete pods.
Here number of replicas is set to 1.



Kubernetes : How to use - Example

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-765bf4c7b4-vvdx	1/1	Running	0	10m

1 – list of pods on the cluster

```
$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
kubernetes-bootcamp-765bf4c7b4	1	1	1	10m

2 – list of « *replicaset* »

```
$ kubectl scale deployments/kubernetes-bootcamp --replicas=4
```

3 – increase desired number to 4

```
deployment.apps/kubernetes-bootcamp scaled
```




Kubernetes : How to use - Example

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-765bf4c7b4-v4ndx	1/1	Running	0	10m

```
$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
kubernetes-bootcamp-765bf4c7b4	1	1	1	10m

```
$ kubectl scale deployments/kubernetes-bootcamp --replicas=4
```

```
deployment.apps/kubernetes-bootcamp scaled
```

```
$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
kubernetes-bootcamp-765bf4c7b4	4	4	4	10m

1 – list of pods on the cluster

2 – list of « *replicaset* »

3 – increase desired number to 4

4 – now 4 replicas are desired
(& ready)



Kubernetes : How to use - Example

```
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-765bf4c7b4-v6w9x 1/1     Running   0           10m

$ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
kubernetes-bootcamp-765bf4c7b4      1         1         1       10m

$ kubectl scale deployments/kubernetes-bootcamp --replicas=4
deployment.apps/kubernetes-bootcamp scaled

$ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
kubernetes-bootcamp-765bf4c7b4      4         4         4       10m

$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-765bf4c7b4-jz3m4 1/1     Running   0           16s
kubernetes-bootcamp-765bf4c7b4-v6w9x 1/1     Running   0           10m
kubernetes-bootcamp-765bf4c7b4-vu89f 1/1     Running   0           16s
kubernetes-bootcamp-765bf4c7b4-x8f9c 1/1     Running   0           16s

$ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE             NOMINATED NODE   READINESS GATES
kubernetes-bootcamp-765bf4c7b4-jz3m4 1/1     Running   0           24s   172.18.0.7      minikube         <none>            <none>
kubernetes-bootcamp-765bf4c7b4-v6w9x 1/1     Running   0           10m   172.18.0.1      minikube         <none>            <none>
kubernetes-bootcamp-765bf4c7b4-vu89f 1/1     Running   0           14s   172.18.0.8      minikube         <none>            <none>
kubernetes-bootcamp-765bf4c7b4-x8f9c 1/1     Running   0           24s   172.18.0.9      minikube         <none>            <none>
```

1 – list of pods on the cluster

2 – list of « *replicasets* »

3 – increase desired number to 4

4 – now 4 replicas are desired (& ready)

5 – 4 pods are available with 4 different cluster ip addresses



Kubernetes : How to use - Example

```
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-765bf4c7b4-v4w9x 1/1     Running   0          10m

$ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
kubernetes-bootcamp-765bf4c7b4      1         1         1       10m

$ kubectl scale deployments/kubernetes-bootcamp --replicas=4
deployment.apps/kubernetes-bootcamp scaled

$ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
kubernetes-bootcamp-765bf4c7b4      4         4         4       10m

$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-765bf4c7b4-jz3m4 1/1     Running   0          16s
kubernetes-bootcamp-765bf4c7b4-v6w9x 1/1     Running   0          10m
kubernetes-bootcamp-765bf4c7b4-vu89f 1/1     Running   0          16s
kubernetes-bootcamp-765bf4c7b4-x3w9c 1/1     Running   0          16s

$ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE             NOMINATED NODE   READINESS GATES
kubernetes-bootcamp-765bf4c7b4-jz3m4 1/1     Running   0          24s   172.18.0.7      minikube         <none>            <none>
kubernetes-bootcamp-765bf4c7b4-v4w9x 1/1     Running   0          30m   172.18.0.1      minikube         <none>            <none>
kubernetes-bootcamp-765bf4c7b4-vu89f 1/1     Running   0          24s   172.18.0.8      minikube         <none>            <none>
kubernetes-bootcamp-765bf4c7b4-x3w9c 1/1     Running   0          24s   172.18.0.9      minikube         <none>            <none>

$ kubectl scale deployments/kubernetes-bootcamp --replicas=1
deployment.apps/kubernetes-bootcamp scaled

$ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
kubernetes-bootcamp-765bf4c7b4      1         1         1       11m
```

1 – list of pods on the cluster

2 – list of « *replicasets* »

3 – increase desired number to 4

4 – now 4 replicas are desired
(& ready)

5 – 4 pods are available
with 4 different cluster
ip addresses

6 – decrease list of replicas to 1



Kubernetes : How to use - Example

```
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-765bf4c7b4-v6w9x 1/1     Running   0           10m

$ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
kubernetes-bootcamp-765bf4c7b4      1         1         1       10m

$ kubectl scale deployments/kubernetes-bootcamp --replicas=4
deployment.apps/kubernetes-bootcamp scaled

$ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
kubernetes-bootcamp-765bf4c7b4      4         4         4       10m

$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-765bf4c7b4-jz3m4 1/1     Running   0           16s
kubernetes-bootcamp-765bf4c7b4-v6w9x 1/1     Running   0           10m
kubernetes-bootcamp-765bf4c7b4-vu89f 1/1     Running   0           16s
kubernetes-bootcamp-765bf4c7b4-x3w9c 1/1     Running   0           16s

$ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE             NOMINATED NODE   READINESS GATES
kubernetes-bootcamp-765bf4c7b4-jz3m4 1/1     Running   0           24s   172.18.0.7      minikube         <none>            <none>
kubernetes-bootcamp-765bf4c7b4-v6w9x 1/1     Running   0           30m   172.18.0.1      minikube         <none>            <none>
kubernetes-bootcamp-765bf4c7b4-vu89f 1/1     Running   0           24s   172.18.0.8      minikube         <none>            <none>
kubernetes-bootcamp-765bf4c7b4-x3w9c 1/1     Running   0           24s   172.18.0.9      minikube         <none>            <none>

$ kubectl scale deployments/kubernetes-bootcamp --replicas=1
deployment.apps/kubernetes-bootcamp scaled

$ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
kubernetes-bootcamp-765bf4c7b4      1         1         1       11m

$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-765bf4c7b4-v6w9x 1/1     Running   0           12m

$
```

1 – list of pods on the cluster

2 – list of « *replicasets* »

3 – increase desired number to 4

4 – now 4 replicas are desired (& ready)

5 – 4 pods are available with 4 different cluster ip addresses

6 – decrease list of replicas to 1

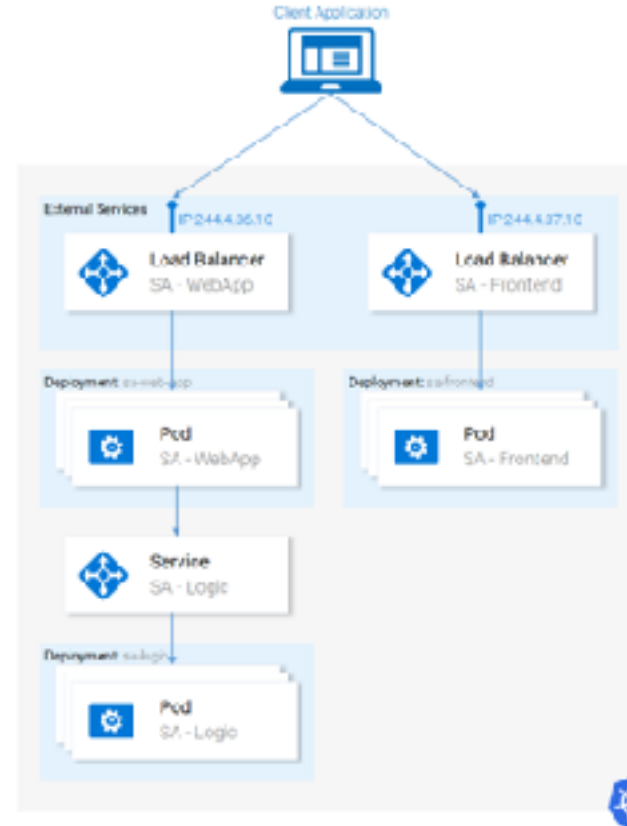
7 – 1 pod is left

Kubernetes : Other features



- Services : Load balancing
- Deployments : Zero-Downtime (BGD : Blue Green Deployment) etc.
- Performance monitoring
- Use of yaml files to fully describes the way applications are deployed (**declarative configuration**)

« kubectl apply -f /path/to/file.yaml »



Kubernetes : Declarative configuration principle



How to : deploy 3 replicas of a piece of software when only 1 is running

Declarative

« Set replica=3 »

Imperative

« Add 2 replicas »

Advantages

Rollback scenario :

- Apply config v2
- * *Bad things happened* *
- Re-apply config v1

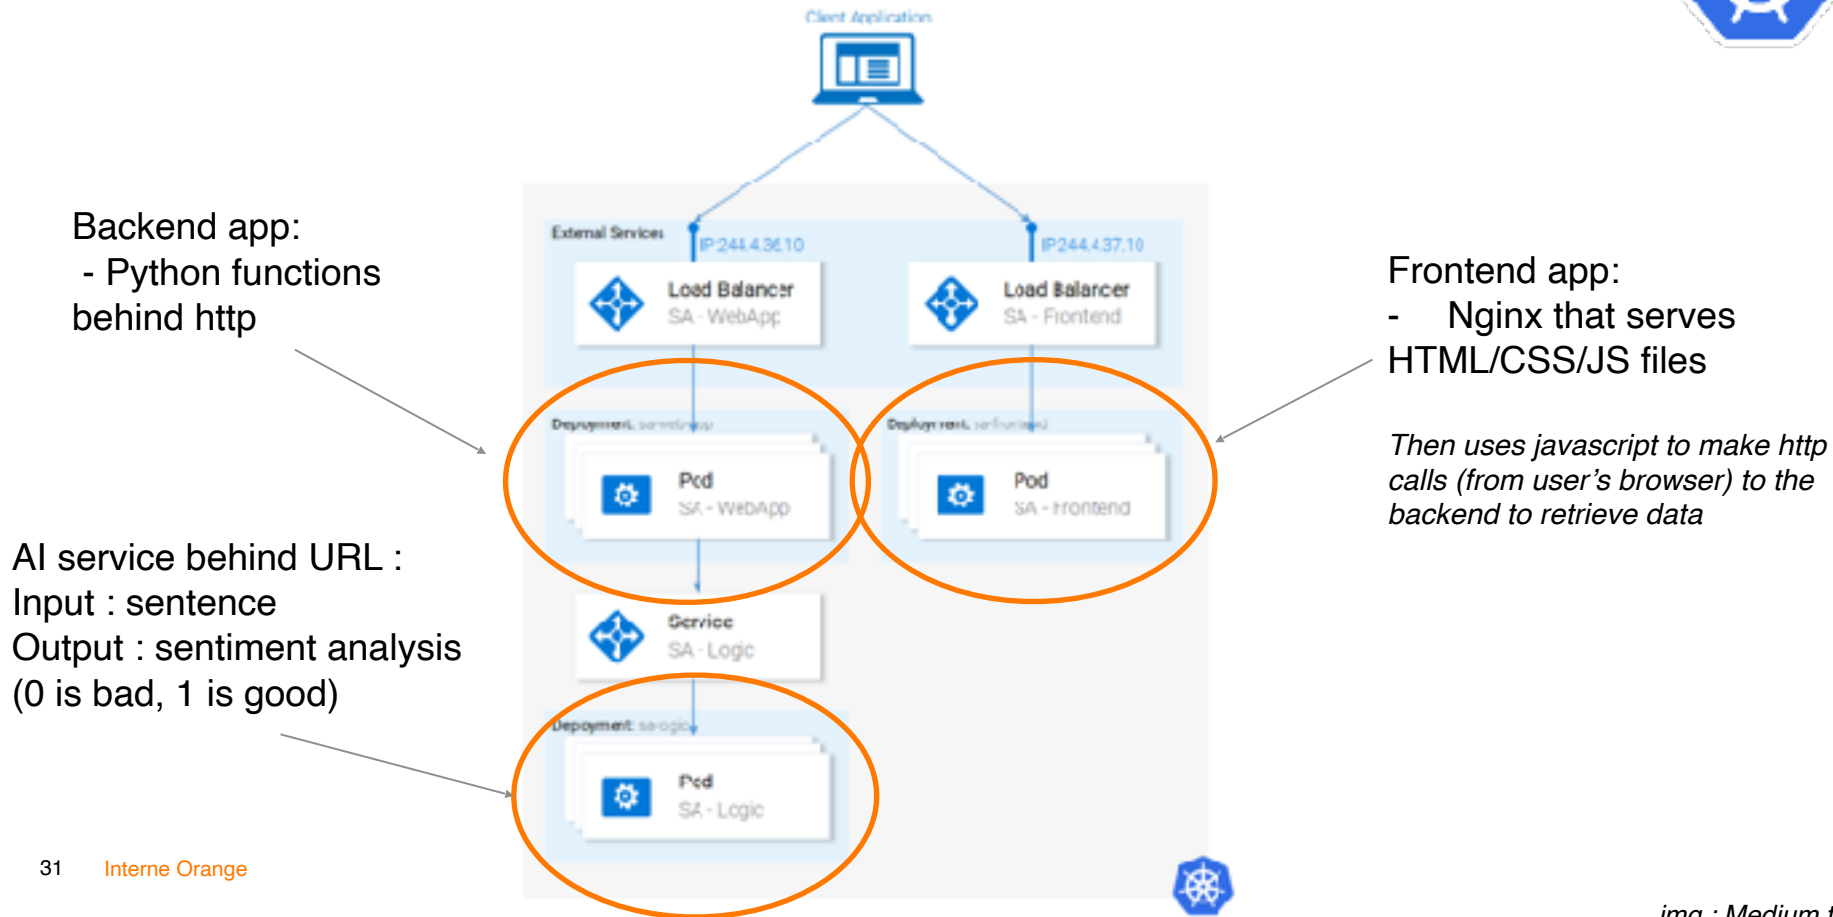
Self-healing algorithm :

- Do current state match desired state ?
 - Yes : re-check
 - No : re-apply desired state

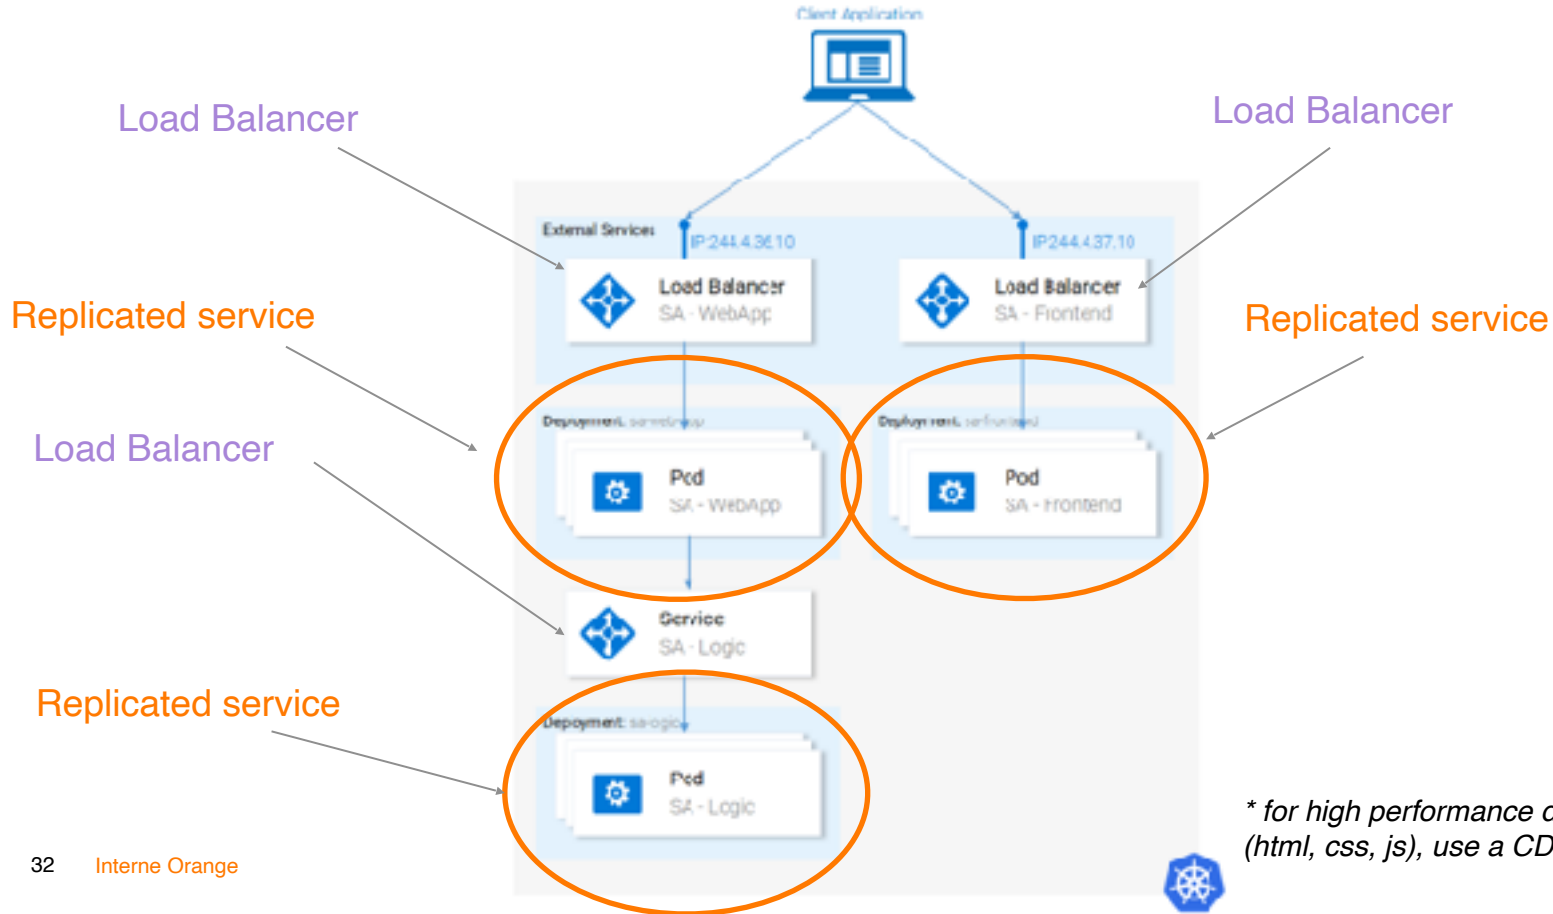
Controllers will check that the current state is identical to the desired state.

Nb: If you manually creates a 4th replica, kubernetes will destroy it.

Kubernetes : Microservices

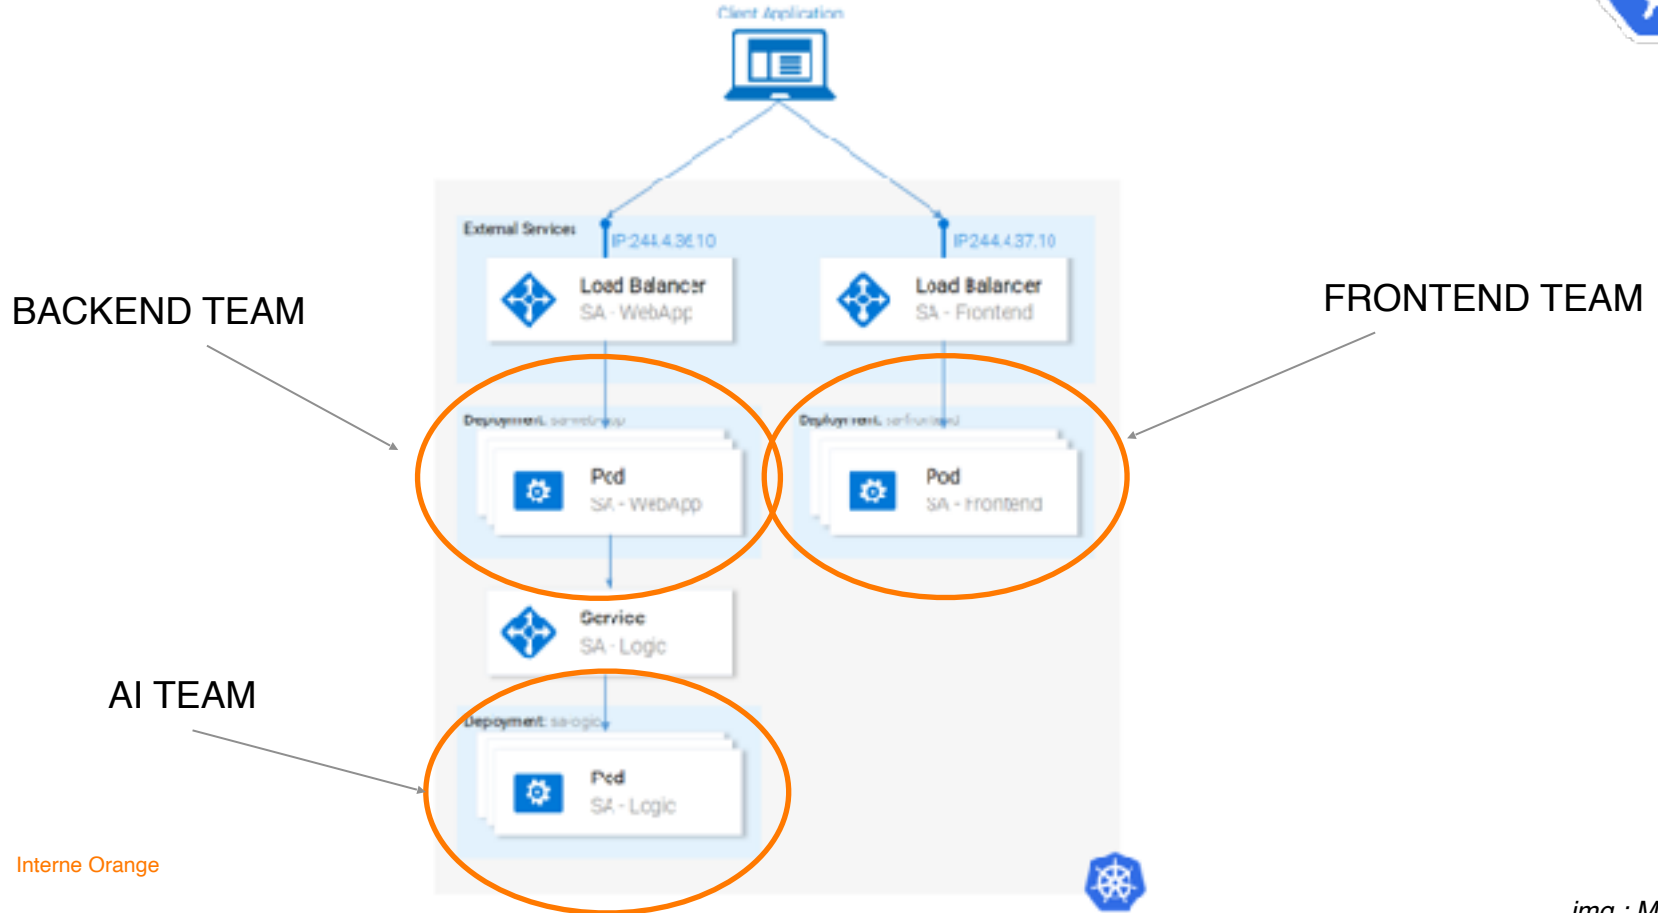


Kubernetes : Microservices workload management



* for high performance delivery of static content (html, css, js), use a CDN may be more efficient.

Kubernetes : Microservices -> 1 service / team



Kubernetes : Resources



- Nice example on Medium :

<https://medium.com/free-code-camp/learn-kubernetes-in-under-3-hours-a-detailed-guide-to-orchestrating-containers-114ff420e882>



THE TWELVE-FACTOR APP

I. Codebase

One codebase tracked in revision control, many deploys

II. Dependencies

Explicitly declare and isolate dependencies

III. Config

Store config in the environment

IV. Backing services

Treat backing services as attached resources

V. Build, release, run

Stricily separate build and run stages

VI. Processes

Execute the app as one or more stateless processes

VII. Port binding

Export services via port binding

VIII. Concurrency

Scale out via the process model

IX. Disposability

Maximize robustness with fast startup and graceful shutdown

X. Dev/prod parity

Keep development, staging, and production as similar as possible

XI. Logs

Treat logs as event streams

XII. Admin processes

Run admin/management tasks as one-off processes

Merci

